

Peer Analysis Report – Kadane's Maximum Subarray Algorithm

Nurkhan Kuanyshov *Partner (Ruslan)*

Implementation Analyzed: Kadane maximum subarray problem

1. Algorithm Overview

The Kadane algorithm finds a contiguous subarray with the maximum sum. It scans the array once while maintaining:

1. currentSum — best sum of any subarray ending at index i .
2. maxSum — best sum seen so far, plus indices [start..end].

Update rule per element $a[i]$:

- If $\text{currentSum} < 0$, start a new subarray at i ($\text{currentSum} = a[i]$, $\text{tempStart} = i$);
- Else extend the current subarray ($\text{currentSum} += a[i]$);
- If $\text{currentSum} > \text{maxSum}$, update maxSum , $\text{start} = \text{tempStart}$, $\text{end} = i$.

Theoretical background:

- Time: linear, $\Theta(n)$ — single pass.
 - Space: constant, $\Theta(1)$ — a handful of scalars.
 - Competes favorably with DP/partitioning approaches: no auxiliary arrays, no recursion.
-

2. Complexity Analysis

2.1 Time Complexity

- Best (Ω): $\Theta(n)$
- Still must read each element once to confirm the maximum and indices.
- Worst (O): $\Theta(n)$
- One linear pass; no nested loops.
- Average (Θ): $\Theta(n)$
- Same single pass regardless of data distribution.

Mathematical justification:

- Let n = array size.
- Single pass over the array $\rightarrow n$ accesses + up to $2n - 2$ comparisons ($\text{currentSum} < 0$, $\text{currentSum} > \text{maxSum}$).
- Constant number of assignments per element (update currentSum, occasional maxSum/start/end) $\rightarrow \Theta(n)$ assignments overall.
- Total operations $\approx n$ accesses + $\leq 2n$ comparisons + $\Theta(n)$ assignments $\rightarrow \Theta(n)$ time, $\Theta(1)$ space.

Comparison with partner's algorithm:

- Partner's version checked metrics != null inside every counter call and flushed the CSV writer inside the loop → extra branches + I/O overhead.
 - KadaneOptimized keeps the same asymptotic complexity but lowers constants by:
 - caching the tracking flag (final boolean track = metrics != null) and grouping metric bumps,
 - caching length in the loop header (for (int i = 1, n = arr.length; i < n; i++)),
 - caching value (final int ai = arr[i]) to avoid repeated reads,
 - moving flush outside the inner loop in the benchmark.
 - Result: same $\Theta(n)/\Theta(1)$, but fewer branches and calls → measurably faster under identical inputs.
-

2.2 Space Complexity

- Auxiliary: $\Theta(1)$ — variables currentSum, maxSum, start, end, tempStart.
 - In-place: operates directly on the input array; no extra buffers.
-

2.3 Recurrence Relations

- Not applicable — iterative scan, no recursion.
-

3. Code Review

3.1 Inefficiency Detection

- Multiple repeated null checks for metrics inside the tight loop (metrics != null) increase branching.
- Repeated arr.length reads and direct arr[i] accesses can be micro-optimized.
- Frequent flush() in benchmark loop adds I/O overhead unrelated to the algorithm.

3.2 Time Complexity Improvements

1. Track flag: cache final boolean track = (metrics != null); and group metric bumps inside a single if (track) block.
2. Length cache: for (int i = 1, n = arr.length; i < n; i++) to avoid repeated bounds fetch.
3. Value cache: final int ai = arr[i]; (fewer array reads).
4. Benchmark I/O: move fw.flush() out of inner loops.

These do not change $\Theta(n)$ but reduce constant factors in practice.

3.3 Space Complexity Improvements

- Already optimal ($\Theta(1)$). Avoid extra objects/arrays.

3.4 Code Quality

- *Clear separation: Kadane (baseline), KadaneMeasured (instrumented), KadaneOptimized (instrumented + micro-opts).*
 - *Deterministic behavior on edge cases: $\{10\} \rightarrow \{10,0,0\}$; all-negative chooses the largest element; empty/null throws `IllegalArgumentException`.*
 - *Unit tests cover classic, all-negative, single-element, and invalid inputs.*
-

4. Empirical Results

4.1 Performance Measurements

Setup: sizes $n \in \{100, 1\,000, 10\,000, 100\,000\}$, 5 trials each; values uniform in $[-1000, 1000]$; time via `System.nanoTime`.

n	Accesses	Comparisons	Time (ms)
100	100	198	0.038
1,000	1,000	1,998	0.206
10,000	10,000	19,998	1.722
100,000	100,000	199,998	1.896

n	Accesses	Comparisons	Time (ms)
100	100	198	0.031
1,000	1,000	1,998	0.209
10,000	10,000	19,998	1.651
100,000	100,000	199,998	1.865

Observation: Linear growth of Accesses/Comparisons confirms theoretical $\Theta(n)$.

Optimized shows lower time in 3/4 sizes ($\approx 1.23\times$ at $n=100$; $1.04\times$ at $n=10k$; $1.02\times$ at $n=100k$). Differences are modest because both variants instrument metrics; the optimizations mainly cut extra branches and I/O overhead. Also time depends on a CPU used for testing.

4.2 Complexity Verification

- Time vs n : near-linear trend confirms $\Theta(n)$.
 - Array accesses vs n : scales linearly with $\sim n$ accesses. Optimized achieves lower runtime with the same access count by reducing branch checks and grouped metric updates.
-

4.3 Comparison Analysis

- Measured \rightarrow Optimized: Reduced branching and repeated reads \rightarrow lower runtime at the same asymptotics.
- With metrics vs without: Instrumentation adds constant overhead; useful during analysis, disabled in production.

4.4 Optimization Impact

<i>Optimization</i>	<i>Effect</i>
Cached track flag	Fewer branch checks in hot loop
Cache $n = \text{arr.length}$	Slightly fewer bound reads
Cache $ai = \text{arr}[i]$	Fewer array reads per iteration
Single flush at end	Removes I/O noise from timing

5. Conclusion

- Kadane runs in $\Theta(n)$ time with $\Theta(1)$ space — optimal for this problem.
- Micro-optimizations keep the algorithm simple yet reduce constant factors, especially when metrics are enabled.
- Empirical results (time/accesses) align with theory.
- Recommendation: use baseline code in production; enable metrics/optimized variant for analysis and benchmarking.